

S&OP Pre-Work Skills

Five standalone workflows for preparing the evidence, options, and decision asks that executive S&OP should receive.

Download only the files you need. No ZIP, email gate, or account.

Run the skills on governed copies of your own project files.

Each skill embeds setup guidance plus its own input and output templates.

Keep every accountable plan and system decision with a named owner.



1 PDF + 1 MD GUIDE

5 STANDALONE SKILLS

0 BUNDLED DATA FILES

READ THIS FIRST

These skills prepare review artifacts. They do not approve a forecast, portfolio change, supply response, financial commitment, or system-of-record update.

1 - DOWNLOAD MAP

Take only the artifact you need

The website shows every file before download. Begin with this PDF, then choose one skill for the review you need. Use all five only when you want to prepare an end-to-end cycle.

Download	What it contains	Use it for
Start Here PDF / MD	Setup paths, first run, quality rules, controls	Understand the method before prompting
Five skill files	Project setup, method, guardrails, and embedded templates	Run one review or the full chain

First safe run

- Choose one review and download its standalone skill.
- Create governed, read-only copies of the relevant project files; start with a narrow scope when needed.
- Attach the skill and project files to ChatGPT, Claude, Gemini, Codex, or another approved AI workspace.
- Run READINESS first: require a source-to-input map and preserve every unknown or conflicting value.
- Run BUILD only when the minimum evidence is sufficient.
- Review the artifact with the named business owner before it becomes the next review's input.

Suggested first prompt

Use this skill with my attached project files. Start in READINESS mode. Create the required source-to-input map, show which file and field satisfies each embedded input, and list anything missing, stale, conflicting, or unclear. Do not invent values or silently resolve conflicts. Move to BUILD mode only when the minimum evidence is sufficient. Then follow the method, return the embedded output contract, and end with the human decisions still required.

2 - THE REVIEW CHAIN

Five hand-offs before one executive decision

A meeting is not complete because its agenda finished. It is complete when the next review receives evidence it can use without reconstructing the work.

Review	What the skill settles	Required hand-off
01 Prepare Data Readiness	Source versions, actuals, exceptions, assumptions	Trusted exception pack
02 Review Portfolio	Launches, phase-outs, lifecycle decisions	Approved portfolio changes
03 Build Demand Plan	Baseline, evidence, judgment, scenarios	Consensus demand and scenarios
04 Test Supply Feasibility	Capacity, materials, inventory, response levers	Feasible response and options
05 Reconcile Plan	Gaps, trade-offs, value, recommendation	Recommendation and decision asks

EXECUTIVE READINESS

READY = EVIDENCE + OWNER + OPTIONS + ASK. This is a Shetty's Desk operating synthesis, not a quoted industry standard.

What each review must avoid

Data: never choose a preferred source silently when versions conflict.

Portfolio: never turn a roadmap date into an approved planning input.

Demand: never hide judgment and overrides inside one consensus number.

Supply: never assume infinite capacity, material, cash, or response time.

Reconciliation: never approve the plan or invent a financial impact.

3 - RUN ONE SKILL

Give the agent a controlled job

The standalone file contains the workflow and both templates. Keep the prompt short; the skill already carries the method. Your job is to provide governed evidence and challenge the result.

Step	Action	Evidence to keep
1	Choose the review and attach its standalone skill.	Skill filename and version
2	Attach governed copies of the mapped project inputs.	Source owner, version, as-of date
3	Run READINESS; move to BUILD only when evidence is sufficient.	Unknowns and conflicting values
4	Review the output against the hand-off gate.	Owner, next check, option, decision ask
5	Approve, reject, or return the artifact for correction.	Named human decision and timestamp

Replace these fields deliberately

- Cycle, horizon, calendar, planning grain, units, and currency.
- Source names, owners, versions, timestamps, and approval states.
- Thresholds, service priorities, constraints, and response windows.
- Decision rights, due points, escalation rules, and system boundaries.

DATA SAFETY

Remove or mask personal, customer-identifying, commercially sensitive, and credential data unless the AI environment and use are explicitly approved.

4 - QUALITY AND CONTROL

Reject polish without traceability

Check	Pass condition	Reject when
Evidence	Every load-bearing number names a source and date	The output invents or silently selects evidence
Ownership	Every exception and assumption has a named owner	The next team must rediscover accountability
Options	Choices use comparable service, capacity, cash, and risk dimensions	One preferred answer hides its trade-off
Ask	The decision, owner, and decision window are explicit	Executives receive a problem with no choice
Boundary	Human approvals and system actions remain visible	The agent appears to authorize the plan

Human decision boundary

- Data owners approve the trusted source and version.
- Portfolio owners approve lifecycle changes.
- Commercial and demand owners approve overrides and consensus demand.
- Operations, supply, procurement, and finance approve response options and values.
- Executives choose the scenario, allocate resources, and commit one plan.
- Authorized system owners make every system-of-record change.

FINAL TEST	Look at the last executive S&OP discussion that stalled. Was evidence missing, ownership unclear, the options incomparable, or the decision ask undefined? Start with the skill that repairs that hand-off.
------------	---